

DAE Mini Assignment

March 30, 2020

1 DAE Mini Assignment

1.1 Analysing the relationship between the production of coal and the amount of electricity generated in South Africa

Many factors influence the amount of electricity that can be generated in a country. Some of these include the number of available generators, the price of resources, such as coal and oil, and the amount of coal and oil produced. The purpose of this notebook is to answer the question: “Does coal production significantly effect the amount of power produced in South Africa?” To do this, the data will go through an analysis and exploration to determine if there is a relationship and, if so, what the relationship is. A logistic regression function will then be implemented to attempt to fit the model and predict future data.

1.1.1 The Data

The first dataset that will be used in this notebook contains information on the production and sales of the mining industry. This information was collected (up until December 2019) by surveying the Department of Mineral Resources and Energy (DMRE) and was published on 13 February 2020. The dataset can be found [here](#) (downloaded on 21 February 2020). In the 44x214 table, the physical volume of mining production (actual indexes and seasonally adjusted indexes), as well as mineral sales for all the resources South Africa mines, can be found. This helps in answering the question as we can use the actual physical volume of mining production index for coal.

The second dataset contains data on electricity generated and available for distribution. The data was collected by statsSA (until December 2019) through surveying establishments in the electricity industry. This dataset can be found [here](#) (downloaded on 21 February 2020). The 24x251 table contains information ranging from the total available electricity for distribution in South Africa to the amount of electricity distributed to each province. To answer the question at hand, we will be looking at the total available for all of South Africa.

The quality of the data is quite high as it is valid, complete, consistent, uniform and accurate as it is taken from a governmental site.

```
[11]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from matplotlib.dates import DateFormatter, AutoDateLocator
from pandas.plotting import register_matplotlib_converters
```

```
from scipy import signal
import seaborn as sns
pd.set_option('display.notebook_repr_html', True)
```

1.1.2 Load Data

```
[19]: #Data for Electricity
power_series = pd.read_excel('Electricity from 2000.xlsx' )
display(power_series.head())
```

	H01		H02	H03	\
0	P4141	Electricity generated and available for distri...		ELEKTS10	
1	P4141	Electricity generated and available for distri...		ELEKIN11	
2	P4141	Electricity generated and available for distri...		ELEKIS11	
3	P4141	Electricity generated and available for distri...		ELEKTR11	
4	P4141	Electricity generated and available for distri...		ELEKTR13	

		H04	\
0		Total - All producers	
1	Physical volume of electricity production		
2	Physical volume of electricity production		
3		Total - All producers	
4		Total - All producers	

		H05	H13	H14	\
0	Electricity available for distribution in Sout...	NaN	NaN		
1		NaN	NaN	NaN	
2		NaN	NaN	NaN	
3		Electricity produced	NaN	NaN	
4	Purchased outside South Africa (import)		NaN	NaN	

	H16	H17	H18	...	M0032019	\
0	Seasonally adjusted	Gigawatt-hours	NaN	...	18751.0	
1	NaN	Index	Base: 2015=100	...	100.4	
2	Seasonally adjusted	Index	Base: 2015=100	...	99.7	
3	NaN	Gigawatt-hours	NaN	...	20943.0	
4	NaN	Gigawatt-hours	NaN	...	707.0	

	M0042019	M0052019	M0062019	M0072019	M0082019	M0092019	M0102019	\
0	19309.0	19229.0	19073.0	18961.0	18463.0	19066.0	18961.0	
1	99.4	105.9	105.2	108.1	103.0	99.6	103.4	
2	102.9	102.2	102.6	101.2	99.2	101.5	100.4	
3	20733.0	22090.0	21947.0	22552.0	21500.0	20781.0	21571.0	
4	689.0	888.0	692.0	825.0	766.0	817.0	900.0	

	M0112019	M0122019
--	----------	----------

```

0  18664.0  18419.0
1    99.0   94.1
2    99.3   97.7
3  20653.0  19640.0
4    903.0   953.0

```

[5 rows x 251 columns]

```

[13]: #Data for Mineral Production
coal_series = pd.read_excel('Coal from 2003.xlsx')
display(coal_series.head())

```

```

      H01                H02      H03 \
0 P2041 Mining: Production and sales FMP20000
1 P2041 Mining: Production and sales FMP20001
2 P2041 Mining: Production and sales FMP21000
3 P2041 Mining: Production and sales FMP23010
4 P2041 Mining: Production and sales FMP23020

```

```

                H04                H05 H06 \
0 Physical volume of mining production Total, gold included NaN
1 Physical volume of mining production Total, gold excluded NaN
2 Physical volume of mining production Coal NaN
3 Physical volume of mining production Iron ore NaN
4 Physical volume of mining production Chromium NaN

```

```

      H16  H17      H18      H25  ... M0032019 M0042019 \
0 Actual indices Index 2015=100 Monthly ... 97.0 90.6
1 Actual indices Index 2015=100 Monthly ... 103.3 95.9
2 Actual indices Index 2015=100 Monthly ... 104.8 95.5
3 Actual indices Index 2015=100 Monthly ... 93.8 90.7
4 Actual indices Index 2015=100 Monthly ... 110.9 112.4

```

```

      M0052019 M0062019 M0072019 M0082019 M0092019 M0102019 M0112019 \
0 103.1 107.5 98.5 102.5 101.9 103.7 101.8
1 110.2 113.6 103.1 106.6 105.8 107.8 105.8
2 112.0 105.7 112.2 109.1 101.3 107.3 102.5
3 105.6 122.0 98.5 95.1 109.9 97.7 107.5
4 119.0 120.6 116.6 117.8 114.0 124.2 121.7

```

```

      M0122019
0 94.2
1 97.6
2 77.6
3 117.4
4 105.3

```

[5 rows x 214 columns]

It is clear to see that the data is very untidy and almost unreadable in its current state. To fix this, both tables are transposed so that the dates can be the rows, and then the data frame is spliced so that they only have the information we require. In this case, we need the total electricity available for distribution from the power table and the actual physical volume of mining production from the coal table. The labels for the months are renamed for readability and then converted to a datetime type and the numerical data is converted to float64.

```
[14]: power_df = pd.DataFrame(power_series.T[11:][0].rename('Total Electricity for_
↳Distribution (GWh)'))
power_df.index.name = 'Month'
power_df.index = power_df.index.map(lambda s: s[4:] + '-' + s[2:4] + '-01')

power_df.index = pd.to_datetime(power_df.index)
power_df['Total Electricity for Distribution (GWh)'] = power_df['Total_
↳Electricity for Distribution (GWh)'].astype(float)

display(power_df.head())
display(power_df.info())
```

```

      Total Electricity for Distribution (GWh)
Month
2000-01-01      15916.0
2000-02-01      15981.0
2000-03-01      16106.0
2000-04-01      16347.0
2000-05-01      16329.0

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 240 entries, 2000-01-01 to 2019-12-01
Data columns (total 1 columns):
Total Electricity for Distribution (GWh)    240 non-null float64
dtypes: float64(1)
memory usage: 3.8 KB

None
```

```
[15]: coal_df = pd.DataFrame(coal_series.T[10:][2].rename('Actual Coal Index'))
coal_df.index.name = 'Month'
coal_df.index = coal_df.index.map(lambda s: s[4:] + '-' + s[2:4] + '-01')
coal_df.index = pd.to_datetime(coal_df.index)
coal_df['Actual Coal Index'] = coal_df['Actual Coal Index'].astype(float)
display(coal_df.head())
display(coal_df.info())
```

```

      Actual Coal Index
```

Month	
2003-01-01	86.5
2003-02-01	82.9
2003-03-01	87.9
2003-04-01	89.7
2003-05-01	101.8

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 204 entries, 2003-01-01 to 2019-12-01
Data columns (total 1 columns):
Actual Coal Index    204 non-null float64
dtypes: float64(1)
memory usage: 3.2 KB
None
```

The data is now much easier to read compared to the raw data and can now be analysed.

```
[16]: #Making both data sets equal size.
#Coal starting 1 month before electricity and electricity ending 1 month after
if np.shape(power_df) != np.shape(coal_df):
    power_df = power_df[37:]
    coal_df = coal_df[:-1]
coal_temp = coal_df.copy()
coal_temp.index = power_df.index
```

For this notebook, we will be looking at the previous month's coal production and comparing it to the current month's electricity generated.

1.1.3 Exploratory Analysis

```
[17]: #Calculating the correlation coefficient
correlcoeff = np.corrcoef(power_df['Total Electricity for Distribution_
    ↳(GWh)'], coal_df['Actual Coal Index'])[0,1]
print('The correlation coefficient of the data is ', correlcoeff )
```

The correlation coefficient of the data is 0.23235245004885707

The correlation coefficient shows that the production of coal and the power generated for distribution in South Africa have a weak positive relationship. Even though the relationship is weak, it still shows that there is a correlation and that there is cause for comparing the datas and answering the question.

```
[18]: #Plotting the two data sets on the same set of axes
%matplotlib notebook
register_matplotlib_converters()
fig,ax = plt.subplots()
```

```

ax2 = ax.twinx()

locator = AutoDateLocator()
formatter = DateFormatter('%Y')

ax.xaxis.set_major_locator(locator)
ax2.xaxis.set_major_locator(locator)

ax.xaxis.set_major_formatter(formatter)
ax2.xaxis.set_major_formatter(formatter)

ax.plot(power_df.index, power_df['Total Electricity for Distribution (GWh)'],
        color = 'red')
ax.set_ylabel('Total Electricity for Distribution (GWh)', fontsize = 14, color =
        'red')
ax.set_xlabel('Years')

ax2.plot(coal_temp.index, coal_temp['Actual Coal Index'], color = 'blue')
ax2.set_ylabel('Actual Coal Index', fontsize = 14, color = 'blue')

sns.set()
ax.xaxis_date()
ax2.xaxis_date()
plt.show()

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Plotting the two data sets in figure 1, visually shows us that they are correlated since both plots are quite similar. The graph also shows us that both data sets experience regular fluctuations every calendar year. This is known as seasonality and may be the cause for the low correlation coefficient of the two data sets. Using the cross-correlation graph and the autocorrelation graph will help in identifying the lags in the data and to detect the seasonality. The data sets will be adjusted for the seasonality and then re-analysed and explored.

[]: